

Latency Optimization of Cascaded Voice Dialogue Systems with a Pipeline-Parallel Streaming Architecture

1st Haihua Mo

School of Computer, Electronics
and Information, Guangxi University
Nanning, China
hua0424@163.com

2nd Zhengyou Liang

School of Computer, Electronics
and Information, Guangxi University
Nanning, China
zhyliang@gxu.edu.cn

Abstract—Cascaded ASR–LLM–TTS dialogue systems preserve modularity but leave speech recognition and prompt prefilling on the post-utterance critical path. We present a pipeline-parallel architecture that overlaps chunked Whisper recognition with incremental LLM prefilling. Silero VAD and a timestamp-aligned sliding window produce an append-only downstream transcript stream, while a reusable key–value cache moves most prompt processing into the user’s speaking time. We evaluate 1,132 synthesized MultiWOZ/CrossWOZ utterances, concatenated human-read speech from LibriSpeech and AISHELL-1, a matched LocalAgreement-2-style baseline, and a unified time-to-first-audio (TTFA) path. On the same 498-sample synthetic subset, the overall TTFT reduction is 74.3% on the original platform and 70.4% on a second dual-RTX 3090 platform. For Extra Long human-read speech, reductions are 67.5% on LibriSpeech and 65.7% on AISHELL-1. Streaming ASR accounts for most of the measured gain; incremental prefilling has a smaller independent effect under Qwen2-7B. Unified-timeline measurements reduce median TTFA from 22.27 s to 3.11 s, although TTS remains a variable bottleneck. Internal ASR hypotheses can still drift after commitment, but no downstream rollback is issued; an exploratory 50-sample semantic evaluation reports a mean response-embedding cosine of 0.883.

Index Terms—intelligent systems, human–computer interaction, voice dialogue system, streaming architecture, pipeline parallelism, incremental prefilling, end-to-end latency

I. INTRODUCTION

Spoken interfaces are used in service robots, intelligent transportation, and smart assistants. This work considers a conventional modular cascade composed of automatic speech recognition (ASR), a large language model (LLM), and text-to-speech (TTS). In the serial configuration studied here, downstream understanding and generation begin only after endpoint detection and full-utterance recognition, leaving recognition and prompt processing on the post-speech critical path. The resulting waiting time motivates the pipeline-parallel schedule evaluated in this paper.

Large-model voice interfaces improve the range of dialogue tasks that can be addressed but also introduce substantial inference cost. OpenAI reports that GPT-4o can respond

to audio input in as little as 232 ms, with an average of 320 ms [1]. These vendor-reported latency figures motivate treating response latency, alongside recognition and response quality, as an explicit evaluation dimension.

End-to-end multimodal models and modular cascades present different engineering trade-offs. The system studied here decomposes processing into independently configured VAD, ASR, LLM, and TTS modules connected in sequence, allowing the scheduling contribution to be evaluated without changing the underlying Whisper or Qwen2 weights. In the serial baseline, total system latency depends on both component processing times and the order in which data become available. For analysis, we express this serial latency budget as a simple sum:

$$T_{\text{total}} = T_{\text{VAD}} + T_{\text{ASR}} + T_{\text{Prefill}} + T_{\text{Decode}} + T_{\text{TTS}} + T_{\text{Net}} \quad (1)$$

Here T_{VAD} denotes voice-activity and turn-boundary processing, T_{ASR} speech recognition, T_{Prefill} prompt prefilling, T_{Decode} autoregressive decoding, T_{TTS} speech synthesis, and T_{Net} communication and serialization. Component overlap means that (1) is a conceptual serial budget rather than an identity for the proposed pipeline. We therefore report two timestamp-defined metrics. Time to first token is

$$\text{TTFT} = t_{\text{first token}} - t_{\text{speech end}}, \quad (2)$$

and time to first audio is

$$\text{TTFA} = t_{\text{first playable PCM}} - t_{\text{speech end}}. \quad (3)$$

For the prerecorded legacy experiments, speech end is operationalized as completion of feeding the waveform. TTFT includes pipeline closure, remaining ASR and prefill work, and first-token selection, but excludes TTS. TTFA additionally includes response-text readiness, the local TTS request, and initial synthesis until at least 1324 bytes of non-empty playable PCM are available. The prototype uses in-process queues for ASR/LLM communication and a localhost TTS service, so it does not measure a distributed network path.

In conventional serial cascades, several seconds may elapse between speech end and spoken output. For long-form input, full-audio ASR and one-shot LLM prefilling both remain after speech end. As spoken interfaces enter service robots and conversational assistants, reducing this residual critical path is important for usability. The question addressed here is whether a modular cascade can move recognition and prefill work into the user’s speaking interval without changing the underlying ASR and LLM weights.

II. RELATED WORK

A. Streaming Progress in Automatic Speech Recognition

The processing speed of the ASR system determines how quickly user speech becomes text the LLM can consume. Early streaming ASR architectures prominently included the RNN-Transducer (RNN-T) [2]. Its recurrent structure supports incremental audio processing, although sequential dependencies limit parallelism over time. Transformer self-attention instead permits broad contextual interaction [3]; without modifications that restrict the available context or attention pattern, however, the vanilla formulation does not provide bounded-latency streaming inference.

Gulati et al. proposed Conformer, which combines convolutional local feature extraction with Transformer-based global modeling [4]. Streaming adaptations commonly constrain recognition to chunks or stabilize outputs across repeated partial hypotheses. For Whisper, the LocalAgreement policy provides one such output-stabilization approach [5], [6]. Whisper itself is an encoder–decoder model developed primarily for offline transcription [7], so applying it incrementally requires explicit context and commitment policies.

B. LLM Inference Acceleration

Once ASR transcription completes, LLM inference can become another major latency bottleneck because autoregressive generation repeatedly accesses an expanding context and its cached state [8]. KV caching stores historical Key/Value states so that they need not be recomputed at every decoding step. vLLM applies PagedAttention to manage this cache for serving workloads [9]. For prompt processing, FlashAttention restructures exact attention around tiled memory access and reports improved throughput in evaluated settings [10]. StreamingLLM retains initial attention-sink tokens together with a recent-token window, enabling bounded-cache streaming over long sequences without claiming retention of arbitrary distant content [11].

C. Full-Pipeline Optimization of Voice Dialogue Systems

Beyond optimizing ASR or the LLM in isolation, coordinated scheduling across the whole pipeline also matters for overall latency. In speech processing tasks such as speech translation, cascaded designs reuse separate components but may incur intermediate-representation loss, error propagation, and accumulated inter-module latency [12]. Recent speech-native dialogue systems instead model or generate speech within an integrated architecture [13]. Representative systems include Mini-Omni [14], Moshi [15], and LLaMA-Omni [16].

D. Contributions of This Work

Much prior work focuses on ASR-side streaming or on speech-native multimodal generation. Complementary to those directions, this paper proposes a fine-grained pipeline-parallel design (hereafter System B) that schedules an existing non-streaming ASR model and text LLM incrementally without changing their weights. Our main contributions are:

- **Streaming ASR context management.** An adaptive Whisper sliding window combines Silero VAD, dynamic buffering, prefix–suffix acoustic context, and timestamp-aligned commitment to produce an append-only downstream text stream.
- **Incremental LLM prefilling.** Newly committed fragments are prefetched into the LLM’s KV cache during speech. Historical cached states are reused, relocating prompt computation into the input interval rather than claiming a lower asymptotic complexity.
- **Broader latency and robustness evaluation.** In addition to the synthesized MultiWOZ/CrossWOZ benchmark, we evaluate concatenated LibriSpeech/AISHELL-1 human-read speech, noise and speed perturbations, repeated measurements, a matched LocalAgreement-2-style baseline [5], [6], and first-audible-output latency.
- **State-correctness and downstream analysis.** We measure internal ASR drift, tokenizer seams, and response-level semantic consistency, distinguishing the append-only downstream contract from stronger and unsupported claims of internal hypothesis or KV-state equivalence.

III. PRINCIPLES AND METHODS

A. Automatic Speech Recognition

1) *Whisper Architecture and Its Offline Nature:* We adopt OpenAI’s Whisper as the speech-to-text module. Whisper is an end-to-end Transformer encoder–decoder model trained on large-scale weakly supervised speech–text data. Its primary report describes strong zero-shot performance across multiple languages and evaluated distribution-shift and noise benchmarks [7].

As shown in Fig. 1, the Whisper large-v3-turbo configuration uses 128-dimensional log-Mel features; this value is checkpoint-specific, whereas the original Whisper models used 80 Mel bins [17]. A convolutional front end downsamples these features before the Transformer encoder, and the decoder attends to the encoded representation while emitting text and timestamp tokens. These timestamps support the alignment policy used by our append-only interface. The archived run metadata identifies the model as `turbo` but does not preserve an immutable repository revision, so the cited public configuration documents the model family rather than the exact artifact downloaded during the experiments.

2) *Streaming Challenges and Our Approach:* Whisper’s encoder can use both past and future frames within an input window. When it is repeatedly applied to overlapping partial audio, words near chunk boundaries may be omitted, repeated, or revised, and the trailing hypothesis can change as more




Fig. 1. Overall architecture of the Whisper model and schematic of feature downsampling

audio arrives [5], [6]. A downstream LLM therefore requires an explicit policy that decides which partial text is stable enough to commit.

We address both problems with segmented input, overlapped context, and controlled commitment. The front end uses VAD to split PCM into speech segments; the ASR side retains acoustic context around the target; and the output side uses word-level timestamps to commit only the current output region. Once delivered, a fragment is not withdrawn, even if later internal recognition changes that region.

B. LLM Inference

1) *Transformer Attention*: Large language models typically adopt a decoder-only Transformer architecture. Within the LLM alone, first-token computation consists mainly of prompt prefill and the first decoding step. The system-level TTFT in (2) additionally includes pipeline closure and the ASR tail. The core of the Transformer is scaled dot-product attention:

$$\text{Attention}(Q, K, V) = \text{softmax}\left(\frac{QK^\top}{\sqrt{d_k}}\right)V \quad (4)$$

Here Q (Query), K (Key), and V (Value) are obtained from the input vectors through linear projection matrices W_Q , W_K , and W_V , respectively. The Key dimension d_k rescales the dot products so that $QK^\top$ does not suffer variance inflation as dimensionality grows, which would saturate the softmax and shrink gradients. This mechanism lets the model aggregate information from any position of the sequence within a single time step, capturing long-range dependencies.

2) *KV Cache: Mathematical Principle and Complexity*: Text generation in an LLM is autoregressive: to generate the N -th token, the model uses the preceding context. Under dense self-attention, a full uncached forward pass over a sequence of length N has an $O(N^2)$ attention cost [3]. If such a full pass were naively repeated at every generation step, the illustrative cumulative cost would be

$$T_{\text{total}}^{\text{naive}} = \sum_{t=1}^N O(t^2) = O\left(\frac{N(N+1)(2N+1)}{6}\right) \approx O(N^3) \quad (5)$$

Modern inference frameworks avoid literally repeating those full forward passes by caching historical Key/Value states [9]. The preceding cubic expression is therefore an intentionally naive reference model, not the complexity of the implementation evaluated here. With cached causal decoding, the model computes new Key/Value states for the current token and appends them to the stored history:

$$K_{\text{cache}}^{(t)} = \text{Concat}\left(K_{\text{cache}}^{(t-1)}, k_t\right) \quad (6)$$

$$V_{\text{cache}}^{(t)} = \text{Concat}\left(V_{\text{cache}}^{(t-1)}, v_t\right) \quad (7)$$

Attention then involves only the interaction of the current query q_t with the cached history K_{cache} and V_{cache} : compute $q_t K_{\text{cache}}^\top$ and take the weighted sum over V_{cache} . The KV cache thus lowers the per-step decoding complexity from $O(t^2)$ to $O(t)$ —still linear in context length—so the cumulative cost of generating N tokens is about $\sum_{t=1}^N O(t) =$

$O(N^2)$, at the price of cache memory that grows linearly with sequence length.

The KV cache plays two distinct roles here. During autoregressive decoding, both systems reuse historical Key/Value states. System B additionally applies cache reuse to prompt arrival: whenever a committed transcript fragment arrives, only its token sequence is forwarded and appended to the cached prompt state. This does not eliminate attention computation or change the asymptotic order of processing the complete input. It relocates prompt work into the speaking interval and thereby reduces the computation remaining after speech end. Separately tokenizing fragments is not guaranteed to produce the same token sequence as tokenizing their concatenation once; therefore, the implementation claims append-only downstream state, not token-sequence, KV-state, or output-distribution identity with one-shot preprocessing. Fig. 2 contrasts one-shot prompt prefill in System A with incremental prompt prefill in System B.

C. Evaluation Metrics

1) *Time to First Token and Time to First Audio*: TTFT follows (2). In legacy prerecorded runs, $t_{\text{speech end}}$ is the waveform feed-completion timestamp; in the unified TTFA run, all timestamps share one `perf_counter_ns` clock. TTFA follows (3), and “playable” means the first non-empty PCM accumulation of at least 1324 bytes. Its closed timeline is speech end $\rightarrow$ feed end $\rightarrow$ pipeline input close $\rightarrow$ first token $\rightarrow$ response text ready $\rightarrow$ TTS request $\rightarrow$ first playable PCM. We define

$$t_{\text{feed_to_close_wait}} = t_{\text{input close}} - t_{\text{feed end}}. \quad (8)$$

This interval includes scheduling before and after explicit flush and must not be interpreted as flush computation alone.

2) *Recognition and Statistical Reporting*: Word error rate (WER), the standard metric for ASR accuracy, is computed from the Levenshtein edit distance:

$$\text{WER} = \frac{S + D + I}{N_{\text{ref}}} \quad (9)$$

where S , D , and I are substitution, deletion, and insertion counts and N_{ref} is the reference length. We primarily report WER for English and CER for Chinese. Table labels distinguish mean-utterance error (the arithmetic mean of sample-level rates) from corpus error $\sum(S+D+I)/\sum N_{\text{ref}}$. Latency tables report the sample mean, sample standard deviation (ddof=1), and tail percentiles. Paired comparisons use 10,000 bootstrap resamples with seed 20260821 for percentile 95% confidence intervals and two-sided Wilcoxon signed-rank tests; Holm correction is applied within the duration-group and augmentation families. Fig. 6 shows observed P5–P95 bands, not confidence intervals.

IV. SYSTEM DESIGN

A. Overall Architecture

1) *Objective and Logical Topology*: The core design objective is to shift from the traditional “receive everything,

then process everything” paradigm to an “incremental receive, streaming process” paradigm, driving first-token latency toward its theoretical lower bound. The logical topology is shown in Fig. 3, which depicts the streaming parallel data flow built on multi-threaded producer–consumer queues. Audio enters the segmentation module continuously as fixed-length PCM chunks; Silero VAD performs activity detection on an accumulation buffer and emits speech segments [18]. The ASR module concatenates several speech segments, invokes Whisper for transcription, and emits append-only text fragments according to the prefix/suffix context policy. Text fragments are passed through a thread-safe queue to the LLM module, which keeps updating the KV cache and starts generation when the termination marker arrives, thereby moving as much LLM prefill computation as possible forward to overlap with audio input.

The overall architecture is a segment–transcribe–prefill–generate pipeline: as soon as an upstream module produces a new audio or text segment, the downstream module can start computing without waiting for the full input. The system comprises three subsystems: a Streaming Audio Segmenter for voice activity detection and segmentation, a Context-Aware ASR Engine for context concatenation and output stability, and an Incremental LLM Inference Service for incremental prefilling and final generation. The KV cache in the LLM part is a general optimization for Transformer decoder inference; we directly reuse the `past_key_values` mechanism built into the inference framework for state reuse.

2) *Key Engineering Details*: To keep the streaming pipeline real-time and controllable in a single-machine prototype, the implementation splits the chain into a multi-threaded pipeline decoupled by thread-safe queues, with state objects carrying cross-segment context.

For communication, modules do not call each other synchronously; they communicate asynchronously through a producer–consumer model. In the implementation, the audio-chunk queue, audio-segment queue, and text-segment queue are all instances of `queue.Queue`; consumers loop on `get` with a timeout, and a designated end-event marks the end of production, avoiding single-point blocking that could deadlock the pipeline. In the ASR stage, the system must simultaneously accept segmented audio from upstream and stream transcribed text downstream, so it is split into two sub-threads, a collector and a transcriber: the former continuously receives upstream speech segments and writes them to a waiting queue; the latter, when the trigger condition is met, concatenates them in batch, calls Whisper for transcription, and pushes the current commit region to the downstream LLM, decoupling audio collection from model inference.

State management uses a hybrid stateful/stateless design. The segmentation module maintains the accumulated audio buffer, segment boundaries, and indices in `StreamState`. The ASR side manages pending and transcription-window queues in `ASRCache`. The LLM persists `past_key_values`, `attention_mask`, and last-step prefill logits in `KVCache`. Together these objects main-





Fig. 2. Prompt-prefill scheduling. System A prefills the complete prompt after speech end; System B prefills committed fragments incrementally. Both systems use standard KV-cache reuse during autoregressive decoding.



Producer–consumer queues let ASR transcription and LLM prefill overlap with the user’s speech; after End-of-Speech only the last chunk and thread synchronization remain on the critical path.

Fig. 3. Logical topology of the streaming parallel architecture

tain cross-segment state and an append-only downstream interface; they do not guarantee that later internal ASR hypotheses remain unchanged.

B. Streaming ASR Context Management

1) *Strategy*: Streaming ASR exposes a configuration-dependent trade-off between the amount of acoustic context

available to the recognizer and the delay before a fragment can be processed. We therefore evaluate fixed segmentation and recognition-trigger settings within each paired comparison and report the observed behavior, rather than claiming a universal optimum.




Fig. 4. Schematic of the context-aware ASR sliding-window mechanism

2) *Dynamic VAD Segmentation*: The system uses Silero VAD to segment the continuous PCM stream online [18]. Define the input audio stream as a time series $\{x_1, x_2, \dots, x_t\}$. The segmenter maintains an accumulation buffer A_{accum} and, on receiving each new audio chunk c_t , performs $A_{\text{accum}} \leftarrow A_{\text{accum}} \oplus c_t$, where $\oplus$ denotes concatenation. Once the accumulated audio exceeds the detection window, Silero VAD is invoked; a closed speech segment is then emitted to the downstream ASR. The legacy TTFT, ablation, context, human-read, and matched-baseline runs use 500-ms chunks, a Silero threshold of 0.5, a 500-ms minimum speech duration, a 300-ms minimum silence duration, and a separate 2.0-s ASR recognition trigger. The R7 unified-TTFA run instead records a Silero threshold of 0.5, a 250-ms minimum speech duration, a 100-ms minimum silence duration, and 30-ms speech padding, while retaining the 500-ms chunk and 2.0-s recognition trigger. These parameters are fixed within each paired comparison. Because no VAD-parameter ablation was performed, we do not claim that the selected values optimize responsiveness or segmentation stability.

3) *Sliding-Window Context Awareness*: To reduce boundary effects in streaming recognition, the ASR engine uses a sliding window with prefix context. Fig. 4 details the model input at time step t . The window contains frozen prefix audio, the target segment, and an optional suffix buffer. This overlap is intended to reduce boundary instability; it does not guarantee that later internal re-recognition will reproduce the same text.

First, the window is updated. Let Q_t be the audio-segment queue at time t . Newly arrived segments first enter a waiting queue and are then merged in batch into the main queue by the transcription thread, yielding Q_t . Transcription is triggered when the accumulated queue duration reaches the threshold τ_{asr} and the queue length satisfies (10):

$$|Q_{\text{temp}}| \geq N_{\text{prefix}} + 1 + N_{\text{suffix}} \quad (10)$$

where $|Q_{\text{temp}}|$ is the current temporary queue length and N_{prefix} and N_{suffix} are the configured numbers of preceding and following segments. When the condition is met, or when the final speech segment arrives, the ASR model concatenates all segments in the temporary queue and performs one transcription pass to produce the raw transcript. In practice a minimum-recognition threshold is also applied to avoid the frequent calls and unstable output caused by overly short windows.

Second, timestamp-aligned text extraction maps Whisper words to segment spans by word end time and emits only segments inside the current commit region. The first and final rounds receive explicit boundary handling so no trailing text is intentionally withheld. After emission, the window retains configured prefix and uncommitted suffix segments. A larger suffix can reduce some boundary errors but also delays commitment and adds decoding work; Table V shows that its effect is metric- and language-dependent.

4) *Commit Contract and Tokenizer Seams*: The ASR maintains two state layers. Internal hypotheses may be overwritten when retained audio is recognized again. In contrast, a committed downstream fragment is delivered once and never retracted. There is no KV-cache rollback path: corrections observed internally after commitment are logged but are not propagated to the LLM. Thus “append-only” describes the downstream interface, not immutable internal recognition. A tokenizer seam occurs when fragment-wise tokenization differs from one-shot tokenization,

$$\mathcal{T}(f_1) \parallel \mathcal{T}(f_2) \neq \mathcal{T}(f_1 \parallel f_2). \quad (11)$$

The cache stores the actual fragment-wise token sequence. Measurements of internal drift and seam frequency are reported in Section V-F.

C. Incremental KV-Cache Prefilling for the LLM

1) *Strategy*: In System A, the LLM obtains the complete prompt only after ASR finishes transcribing the whole utterance, and performs prefill plus first-token generation entirely after the end-of-speech endpoint. TTFT therefore rises markedly as transcript length grows. Applying that serial paradigm unchanged to a setting where computation could proceed during audio input wastes the time window before the endpoint and squanders the available parallelism.

System B reduces residual computation by retaining the `past_key_values` produced by each committed fragment. A new fragment triggers a forward pass only for its own tokens, after which the updated cache replaces the previous state. When input closes, any remaining fragment and the generation prompt are processed before decoding begins. The amount of work left at closure depends on the final segment and queue state; the design moves most, but not necessarily all, prefill into the speaking interval. Because the cache is append-only, revising an already delivered fragment would require rollback or recomputation, which this prototype does not implement.

2) *Core Algorithm*: The core logic of incremental inference: on the first call, the system performs a one-shot prefill without any cache, establishing the initial KV cache; thereafter, each new ASR text fragment from upstream triggers a call to `cache_prompt()`, which updates incrementally on top of the existing `past_key_values`. When the termination marker is received, the system appends the generation prompt (e.g., “assistant:”, telling the model that the next output token begins the assistant’s reply) to the end of the prompt and reuses the last-step logits from prefilling to

Algorithm 1 Incremental KV cache prefill (StreamLLMInference)

Require: current KV cache $C_{\text{prev}} = (P_{\text{prev}}, \text{Mask}_{\text{prev}}, L_{\text{prev}})$; new text fragment T_{new} ; end flag IsEnd ; generation prompt G ; tokenizer $\mathcal{T}$; LLM $\mathcal{M}$

Ensure: updated cache $C_{\text{new}} = (P_{\text{new}}, \text{Mask}_{\text{new}}, L_{\text{new}})$

- 1: **Step 1: fragment preparation and tokenization**
- 2: **if** $\text{IsEnd} = \text{True}$ **then**
- 3: $T_{\text{new}} \leftarrow \text{Concat}(T_{\text{new}}, G)$ $\triangleright$ append generation prompt
- 4: **end if**
- 5: **if** T_{new} is empty **then**
- 6: **raise exception** $\triangleright$ invalid empty fragment
- 7: **end if**
- 8: $\text{Ids}_{\text{new}} \leftarrow \mathcal{T}(T_{\text{new}}, \text{add_special_tokens} = \text{False})$
- 9: **if** Ids_{new} is empty **then**
- 10: **raise exception** $\triangleright$ no valid token in fragment
- 11: **end if**
- 12: **Step 2: attention mask and position IDs**
- 13: $\text{Mask}_{\text{new}} \leftarrow \text{Concat}(\text{Mask}_{\text{prev}}, \mathbf{1}_{|\text{Ids}_{\text{new}}|})$
- 14: $\text{Pos} \leftarrow [|\text{Mask}_{\text{prev}}|, \dots, |\text{Mask}_{\text{new}}| - 1]$ $\triangleright$ explicit position_ids
- 15: **Step 3: forward pass (incremental update)**
- 16: $\text{Outputs} \leftarrow \mathcal{M}(\text{Ids}_{\text{new}}, \text{Mask}_{\text{new}}, \text{Pos}, P_{\text{prev}}, \text{use_cache} = \text{True})$
 $\triangleright$ forward only on new tokens, reusing cached P_{prev}
- 17: **Step 4: cache and logits update**
- 18: $P_{\text{new}} \leftarrow \text{Outputs.past_key_values}$
- 19: $L_{\text{new}} \leftarrow \text{Outputs.logits[:, -1, :]}$ $\triangleright$
if $\text{IsEnd} = \text{True}$, the first token is decoded from L_{new} without an extra forward pass
- 20: $C_{\text{new}} \leftarrow (P_{\text{new}}, \text{Mask}_{\text{new}}, L_{\text{new}})$
- 21: **return** C_{new}

start decoding directly. To avoid positional-encoding misalignment of incremental tokens, the implementation constructs `position_ids` explicitly and encodes new fragments without adding another generation prompt, keeping cache concatenation controlled. Algorithm 1 gives the incremental update procedure corresponding to this implementation. Fig. 5 visualizes the state changes in GPU memory: the blue region on the left is the cached key-value pairs of the historical context, and the green region on the right is the key-value pairs produced by the new fragment. On each fragment’s arrival, incremental prefilling reuses the historical `past_key_values` and runs the forward pass only over new tokens before updating the cache, avoiding repeated prefills over the full prompt in the incremental setting.

3) *Complexity Analysis:* To locate the source of the redundancy that incremental prefilling removes, consider one update with historical context length N and new-fragment length M . If every incremental arrival triggered a fresh prefill over the whole sequence $x_{1:N+M}$ (no reuse of historical state), the dominant self-attention cost would be about $O((N+M)^2)$.




Fig. 5. Schematic of the LLM KV Cache incremental update mechanism

With the KV cache, the K, V of historical tokens are already stored; the system only computes queries for the new tokens and lets them interact with Key/Value of length $N + M$, at a dominant cost of $O(M(N + M))$. When $M \ll N$ this approximates $O(MN)$, i.e., linear scaling in the new fragment. Note that if all incremental fragments accumulate to a final length L , the cumulative cost of incremental prefilling is

$$\sum_j O(M_j L_j) = O(L^2) \quad (12)$$

where M_j and L_j are the length of the j -th new text fragment and the accumulated length, and L is the full sequence length—the same order as offline one-shot prefilling. The key benefit of our strategy is not a lower complexity order but the relocation of prefill computation into the user’s speaking time, so that the computation remaining after the end-of-speech endpoint consists mainly of the final small increment plus thread-synchronization overhead, effectively lowering TTFT. In addition, our implementation keeps the last-step logits of the prefill phase in the KV cache, so the first token can be decoded without any extra forward pass once the termination marker arrives, further trimming the constant overhead of the first-token phase.

D. Test Data Construction Pipeline

1) *Overview:* Systematically validating the streaming architecture under controlled conditions requires dialogue data of steadily increasing length, so we built an automated data-synthesis pipeline. It converts two standard multi-turn dialogue datasets, MultiWOZ (English) [19] and CrossWOZ (Chinese) [20], into progressively lengthening long-speech test data with precise duration annotations, and then generates audio via TTS.

2) *Cumulative Dialogue Generation:* MultiWOZ and CrossWOZ both contain multi-turn dialogues, but individual user turns are usually short and cannot directly cover long-speech input. To simulate a user stating a complex request in one pass or continuously adding information, we construct long input samples with a cumulative dialogue strategy. As illustrated in Table I, the algorithm walks through the dialogue history turn by turn, concatenating the user-side inputs and

TABLE I
DIALOGUE DATA ACCUMULATION STRATEGY (EXAMPLE TRANSLATED FROM CHINESE)

Turn	User	Agent	Accumulated text
1	Book a flight	Where to?	Book a flight
2	To Beijing, leaving tomorrow	Which cabin class?	Book a flight, Where to?, To Beijing, leaving tomorrow
3	Business class	OK, booking it now	Book a flight, Where to?, To Beijing, leaving tomorrow. Which cabin class? Business class

system replies of the current and preceding turns into input sequences of increasing length.

This strategy generates continuous speech samples from 3 seconds to over 60 seconds grounded in real semantic context, covering interaction forms from short commands to extended statements and remedying the lack of long-speech samples in conventional datasets.

3) *Data Processing Pipeline*: Data construction proceeds in three strictly ordered stages.

The first is history accumulation and filtering. The strategy traverses the source datasets and applies the cumulative generation logic. To focus on long-speech bottlenecks, generated samples are sorted by text length in descending order and the longest N dialogue fragments are selected first, with maximum text-length thresholds (2,050 characters for English, 720 for Chinese) to prevent out-of-memory errors; an upper bound on speech length also avoids needlessly inflating experiment time.

The second stage is concurrent TTS synthesis. We use a CosyVoice-based HTTP synthesis endpoint to produce controlled synthetic waveforms from the accumulated text [21]. A batch client issues asynchronous requests to the endpoint; the resulting audio remains synthetic and is not treated as human speech. The archived artifacts do not preserve an immutable CosyVoice checkpoint or service revision, so [21] identifies the model family rather than a specific hosted Alibaba service.

The third is duration calibration and metadata synchronization. Because generative TTS speaking rate is nondeterministic, estimating duration from character counts is error-prone. The pipeline parses the generated WAV file headers to obtain physical durations in seconds (accurate to milliseconds), which are written back to the test-metadata JSON as ground truth for the X-axis (input duration) of subsequent experiments.

4) *Test-Set Groups*: For fine-grained analysis of latency across durations, the generated samples are divided into five predefined analysis groups by audio duration (Table II).

This grouping is used throughout the experimental analysis in Section V to compare how the architectures handle speech of different lengths.

5) *Concatenated Human-Read Speech and Augmentation*: To test acoustic conditions not represented by TTS, we construct long-form samples from LibriSpeech test-clean [22] and AISHELL-1 [23]. LibriSpeech utterances are concatenated in chapter order and AISHELL-1 utterances in speaker order. Uniformly sampled silences $U(0.2, 1.0)$ s are inserted between source utterances with seed 42. Each language contributes 75 clean samples (30 Long, 30 Very Long, and 15 Extra Long),

adding multiple speakers and recording conditions; AISHELL-1 further broadens Mandarin speaker and accent coverage. For each language, six 30-sample variants add MUSAN [24] noise at SNR 20, 15, and 10 dB, apply speed factors 0.9 and 1.1, or add 15-dB babble. Speed-modified samples are regrouped using their new durations. A five-sample manual listening check passed intelligibility and alignment checks, while confirming that concatenation seams can be audible. We therefore describe these data as concatenated human-read speech rather than spontaneous dialogue.

V. EXPERIMENTS AND ANALYSIS

A. Experimental Protocol

System A is a serial non-streaming cascade: full-audio Whisper large-v3-turbo transcription is followed by one-shot Qwen/Qwen2-7B-Instruct [25] prompt prefill and generation. System B uses the same loaded weights but applies the streaming ASR and incremental prefill mechanisms described above. The archived run metadata records the model identifiers as `turbo` and `Qwen/Qwen2-7B-Instruct`; immutable model-repository revisions were not preserved and are therefore reported as unavailable rather than inferred retrospectively. Tables III–V and Fig. 6 use the original 2025 dual-RTX 3090 platform. Repeated measurements and Tables VI–VIII use a second platform, a KVM virtual machine with an Intel Xeon Gold 6133 at 2.50 GHz and two RTX 3090 GPUs. Absolute latency is platform-dependent, so values from the two platforms are neither scaled nor mixed within a comparison. In both environments, Whisper runs on one GPU and Qwen2-7B-Instruct on the other, batch size is one, and three real-audio warm-up runs precede measurement.

Tables III–V use a 50-token generation cap. The state, semantic, and unified TTFA experiments use 128 tokens. Effective sampling uses temperature 0.1 and top- p 0.9. Although the historical interface accepted `repetition_penalty=1.1`, post-experiment inspection found that the custom decoder did not apply it; all compared arms therefore used no effective repetition penalty. The original-platform mixed synthetic set contains 1,132 samples: 630 from MultiWOZ and 502 from CrossWOZ. Paired exclusions are stated with each experiment.

B. Latency versus Speech Length

System B adds fixed overhead for Short and Medium samples, but its mean TTFT changes little over the three longer groups on the original platform. For Long, Very Long, and Extra Long input, mean reductions are 34.6%, 65.6%, and 83.9%;

TABLE II
AUDIO DURATION GROUPS

Group	Duration (s)	Typical scenario	Experimental focus
Short	$T < 5$	Short commands	Baseline latency
Medium	$5 \leq T < 15$	Multi-intent statements	Everyday dialogue performance
Long	$15 \leq T < 30$	Complex long sentences	Core advantage region of streaming
Very Long	$30 \leq T < 60$	Extended expositions	Tail-latency growth
Extra Long	$T \geq 60$	Extreme stress test	Stability and OOM boundary

TABLE III
TTFT BY DURATION ON THE ORIGINAL PLATFORM (MS)

Group	n	System B			System A		
		Mean $\pm$ std	P95	P99	Mean $\pm$ std	P95	P99
Short	35	648.17 $\pm$ 70.76	744.76	797.17	533.42 $\pm$ 65.44	631.82	674.06
Medium	89	959.53 $\pm$ 184.10	1261.82	1397.39	923.42 $\pm$ 170.86	1196.04	1281.49
Long	121	1126.63 $\pm$ 323.72	1883.55	1978.67	1722.04 $\pm$ 342.63	2310.40	2474.07
Very Long	208	1099.16 $\pm$ 343.51	1768.83	2173.97	3191.97 $\pm$ 697.81	4467.14	4826.85
Extra Long	679	1087.70 $\pm$ 383.16	1717.10	2605.45	6745.57 $\pm$ 2108.79	11313.22	12274.38




Fig. 6. TTFT over 12 equal-frequency duration bins. Lines show bin means and shaded regions show observed P5–P95 ranges, not confidence intervals.

paired bootstrap 95% confidence intervals for the reduction are [30.6%, 38.3%], [63.8%, 67.3%], and [83.3%, 84.4%], respectively, with Holm-adjusted Wilcoxon $p < 2.1 \times 10^{-19}$. The Extra Long absolute reduction is 5657.9 ms (5.66 s). Tail latency is not constant: System B P99 rises from 1978.67 to 2605.45 ms between Long and Extra Long. This $1.32\times$ growth is nevertheless smaller than System A’s $4.96\times$ growth, and the Extra Long P99 is $0.21\times$ the baseline value. The empirical crossover depends on platform and speech distribution; the 15-s group boundary is not asserted as a universal routing threshold.

1) *Repeated Measurements*: On the second platform, 50 fixed Very Long samples were run three times per mode. System B round means were 1434.0, 1482.1, and 1454.9 ms, a range equal to 3.3% of the grand mean; System A means were 4063.3, 4244.2, and 4118.5 ms. With sample standard deviation, System B per-sample coefficient of variation (CV) had mean 5.19%, median 4.05%, P90 10.73%, and maximum 18.96%; the corresponding System A values were 5.23%, 4.65%, 9.92%, and 14.01%. Central behavior was repeatable,

but both modes retained a non-negligible high-variance tail.

C. Ablation and Context Trade-offs

Table IV summarizes the ablation results. The analysis starts from 505 Long-or-longer candidates. Pairwise exclusion removes three runtime errors and four cases in which a streaming arm exceeded 10 s, leaving 498 samples (108 Long, 150 Very Long, and 240 Extra Long). The unfiltered failure/hang rate is therefore 7/505 (1.4%).

Streaming ASR is the dominant source of measured reduction: its gains are 626.13, 2152.87, and 5286.86 ms. Incremental prefill has no monotonic independent effect across the groups: -23.01 , 2.73 , and 40.82 ms relative to Streaming ASR Only. The Extra Long value is 3.3% of the ASR-only TTFT. Thus incremental prefill remains an architectural mechanism for moving prompt work before speech end, but its isolated marginal TTFT effect is small for Qwen2-7B on this platform.

Each context configuration contains 74 MultiWOZ and 76 CrossWOZ samples. Error cells are mean-utterance metrics; the ASR time column pools the streaming and non-streaming records ($n = 300$) and is not interpreted as streaming tail latency alone. Adding one suffix segment reduces CrossWOZ WER/CER from 0.0118/0.0292 to 0.0083/0.0203 while increasing pooled ASR time. Removing all extra context is faster but raises CrossWOZ WER/CER to 0.0210/0.0500. MultiWOZ WER is not monotonic in context size. These data establish a measured trade-off on synthetic speech, not a universal language-specific configuration rule. System A’s near-zero error in this experiment is specific to TTS generated from its reference text.

D. Concatenated Human-Read Speech

On clean LibriSpeech, Extra Long mean TTFT is 4805.1 ± 420.2 ms for System A and 1559.0 ± 490.3 ms for System B, a 67.5% reduction. On clean AISHELL-1 the corresponding values are 5139.9 ± 381.7 and 1762.9 ± 275.4 ms, a 65.7% reduction. The clean human-read System A errors

TABLE IV
ABLATION ON THE ORIGINAL PLATFORM; CELLS ARE MEAN $\pm$ STD [P95/P99] MS

Group	Baseline	Streaming ASR only	Full streaming	ASR gain	KV gain
Long	1690.81 $\pm$ 349.51 [2323.60/2431.03]	1064.67 $\pm$ 307.62 [1691.03/1873.56]	1087.69 $\pm$ 312.46 [1748.08/1859.16]	626.13	-23.01
Very Long	3307.96 $\pm$ 872.57 [4891.94/5468.95]	1155.09 $\pm$ 365.67 [1859.20/2259.80]	1152.36 $\pm$ 363.23 [1828.80/2228.55]	2152.87	2.73
Extra Long	6515.67 $\pm$ 1762.26 [9208.54/9619.23]	1228.82 $\pm$ 411.10 [1981.51/2858.89]	1188.00 $\pm$ 411.07 [1946.51/2812.95]	5286.86	40.82

TABLE V
CONTEXT-WINDOW ACCURACY AND POOLED ASR TIME ON THE ORIGINAL PLATFORM; CELLS SHOW MEAN $\pm$ STD [P95/P99]

Configuration	MultiWOZ WER	MultiWOZ CER	CrossWOZ WER	CrossWOZ CER	Pooled ASR time (ms)
Prefix 1 + suffix 1	.0895 $\pm$.0614 [.1960/.2482]	.0196 $\pm$.0208 [.0611/.0799]	.0083 $\pm$.0160 [.0431/.0659]	.0203 $\pm$.0276 [.0747/.1054]	1327.48 $\pm$ 645.03 [2716.80/3049.16]
Prefix 1 + suffix 0	.0813 $\pm$.0611 [.1882/.2465]	.0221 $\pm$.0218 [.0603/.0842]	.0118 $\pm$.0188 [.0482/.0723]	.0292 $\pm$.0339 [.0927/.1193]	1224.96 $\pm$ 655.85 [2726.57/3018.06]
Prefix 0 + suffix 0	.0619 $\pm$.0678 [.1756/.3086]	.0226 $\pm$.0265 [.0657/.1072]	.0210 $\pm$.0237 [.0633/.0870]	.0500 $\pm$.0374 [.1146/.1316]	1086.16 $\pm$ 690.43 [2641.87/2973.93]

TABLE VI
CONCATENATED HUMAN-READ SPEECH AND AUGMENTATION RESULTS ON THE SECOND PLATFORM

Language	Condition (n)	A TTFT	B TTFT	A error	B error
English	Clean (75)	2805.9	1672.1	.0297 [.0290]	.0565 [.0499]
English	SNR 20 dB (30)	2228.8	1633.1	.0288 [.0270]	.0627 [.0573]
English	SNR 15 dB (30)	2334.3	1801.5	.0255 [.0264]	.0578 [.0556]
English	SNR 10 dB (30)	2201.4	1654.3	.0372 [.0387]	.0817 [.0799]
English	Speed 0.9 $\times$ (30)	2341.8	1602.4	.0458 [.0491]	.0751 [.0751]
English	Speed 1.1 $\times$ (30)	2181.7	1657.0	.0380 [.0440]	.0741 [.0767]
English	Babble (30)	2325.8	3425.1	.0271 [.0296]	.4372 [.4001]
Chinese	Clean (75)	2871.7	1698.8	.1077 [.1118]	.1180 [.1255]
Chinese	SNR 20 dB (30)	2294.6	1696.1	.1224 [.1281]	.1242 [.1299]
Chinese	SNR 15 dB (30)	2290.5	1630.6	.1064 [.1137]	.1110 [.1173]
Chinese	SNR 10 dB (30)	2311.8	1654.7	.1233 [.1275]	.1312 [.1405]
Chinese	Speed 0.9 $\times$ (30)	2236.5	1539.1	.1232 [.1306]	.1380 [.1392]
Chinese	Speed 1.1 $\times$ (30)	2125.4	1623.6	.1510 [.1612]	.1827 [.1918]
Chinese	Babble (30)	2227.7	2082.2	.1360 [.1429]	.2863 [.2662]

TTFT is mean ms. Error is mean-utterance [corpus] WER for English and CER for Chinese.

are no longer near zero: mean-utterance WER is 2.97% on LibriSpeech and CER is 10.77% on AISHELL-1; System B obtains 5.65% and 11.80%. AISHELL-1 references spell numbers in Chinese while Whisper often emits Arabic numerals; this affects 49/75 clean samples (mean CER 0.1476 when affected versus 0.0313 otherwise).

Across the ten SNR and speed conditions, TTFT reduction ranges from 22.8% to 31.6%, with Holm-adjusted paired tests at $p \leq 0.0129$. Babble is a boundary rather than a success case. For English, System B is 47.3% slower in the unpaired overall summary; paired inference on 29 common samples has a confidence interval reaching zero. Chinese babble yields only a 6.5% mean reduction and is not significant ($p = 0.271$). VAD over-triggering creates segment backlog, while Whisper produces empty streaming outputs for 12/30 English and 5/30 Chinese samples. English-babble System B has median 2020.4 ms but mean 3425.1 ms, and the slowest sample takes approximately 21 s after audio completion to drain the ASR queue.

E. Matched Streaming Baseline

We implement a project-internal LocalAgreement-2-style baseline following the partial-hypothesis policy of Liu et al. [5] and its Whisper adaptation [6]. It uses the same Whisper weights, inference engine, and segmentation as Sys-

tem B, with absolute-time-axis commitment, punctuation-robust agreement, sentence-boundary-aware trimming, and a 15-s retained-buffer cap. All three systems are rerun on the same second-platform 498-sample paired subset, and only the final validated baseline implementation is included.

As shown in Table VII, the LA-style baseline is approximately 34% slower than System B (2115.0 versus 1573.9 ms). The paired mean difference is 541.1 ms, with bootstrap 95% CI [485.3, 599.9] and Wilcoxon $p = 1.24 \times 10^{-70}$. Recognition errors remain of the same order, which supports latency attribution but does not establish statistical quality equivalence. Repeated whole-buffer decoding also degrades more on longer input: LA-style TTFT is 2199.7–2230.4 ms for Very Long and Extra Long samples, compared with 1551.1–1637.5 ms for System B.

F. Commitment, Tokenizer Seams, and Semantics

Instrumentation on 50 fixed Very Long synthetic samples records 375 downstream commits and zero rollback deliveries. Internal re-recognition nevertheless changes previously committed-region hypotheses 224 times, affecting 224/425 committed segments (52.7%) and 49/50 samples. Character edit distance has mean 2.3, P90 6, and maximum 16; normalized edit ratio has mean 5.6% and maximum 47.1%. The changes include lexical drift but remain invisible to

TABLE VII
MATCHED BASELINE COMPARISON ON THE SECOND PLATFORM

System	Overall TTFT	P50	P95	Long	Very Long	Extra Long	zh CER	en WER
System A	5310.8±2817.9	4780.9	10462.1	1958.2	3905.9	7697.5	.0674	.1527
System B	1573.9±411.6	1484.4	2296.5	1464.3	1551.1	1637.5	.0732	.1702
LA-2-style	2115.0±693.3	1948.5	3463.2	1741.0	2199.7	2230.4	.0835	.1566

TTFT values are ms; error values are mean-utterance metrics. Overall WER: A .0951, B .1047, LA .1073.

downstream consumers. Concatenated System B text differs from System A full transcription by mean WER 4.93% and CER 2.69%.

Fragment-wise and one-shot token sequences differ at seams in 25/50 samples, with means of 4.36 differing blocks and 5.60 one-shot-side tokens among affected samples. Decoding both token sequences reconstructs identical text in all 50 cases. This supports text-level reconstruction only; it does not establish token, KV-state, or output-distribution equivalence.

Response quality is explored on the same 25 Chinese and 25 English synthetic samples. Replies are independently sampled with the 128-token cap. CLS embeddings from the locally downloaded BAAI/bge-m3 artifact [26], followed by L2 normalization and cosine similarity, yield mean 0.8832 ± 0.0755 (minimum 0.6232, P10 0.7951). The archived artifact was not a versioned Hugging Face snapshot, so its source revision is unavailable. A paired-equivalence judge gives 2.96/5 and 40.0% at or above 4; this measure confounds independent recommendation choices and truncation. A separate blinded intent-satisfaction judge gives System A 3.10/5 and System B 3.04/5; the paired B–A difference is -0.06 with bootstrap 95% CI $[-0.34, 0.22]$. The judge is the Command Code service identifier `deepseek/deepseek-v4-flash` [27], accessed in August 2026 with temperature 0 and deterministic A/B order randomization. No immutable backend build identifier was available. This single-judge, 50-sample analysis is exploratory and is not an equivalence claim.

G. Unified Time to First Audio

Table VIII uses only the R7 unified run on the second platform. Each mode has 25 Chinese and 25 English samples on one monotonic clock. System B sends the first completed sentence to TTS, whereas System A follows its non-streaming full-response text path; this is an end-to-end system-policy comparison, not an isolated TTS comparison with identical text. The local ASR/LLM pipeline uses in-process queues and TTS runs on localhost.

System B reduces median TTFA by 86.0%, from 22269.9 to 3113.7 ms. TTS remains the largest mean component and is highly variable, especially for English. The six components close exactly for all 100 records. The 133.0-ms $t_{\text{feed_to_close_wait}}$ is the complete feed-end-to-input-close interval; independent audit attributes approximately 132.68 ms to waiting before explicit flush, 0.21 ms to explicit flush, and 0.12 ms after flush. It is therefore not a flush-computation cost. On a ten-sample three-round subset, TTFA CV has mean 7.73% and maximum 20.70%. The 128-token cap affects both arms: 84/100 responses terminate at the cap and 16 by EOS.

VI. CONCLUSION AND OUTLOOK

A. Summary

This paper presents a pipeline-parallel cascade that moves chunked ASR and LLM prompt prefilling into the user’s speaking interval. The design preserves the existing Whisper and Qwen2 weights and does not change the asymptotic attention cost; it changes scheduling and state reuse. On the same 498-sample synthetic subset, the overall TTFT reduction is 74.3% on the original platform and 70.4% on the second platform. On the original platform, mean System B TTFT remains near 1.1 s from Long through Extra Long, while P99 grows only $1.32\times$ compared with $4.96\times$ for System A. Concatenated human-read speech reproduces the Extra Long advantage, with reductions of 67.5% on LibriSpeech and 65.7% on AISHELL-1. Ablation shows that streaming ASR supplies most of the measured gain; the independent incremental-prefill effect is small under Qwen2-7B. The matched LA-style baseline is approximately 34% slower than System B. Unified measurement further reduces median TTFA from 22.27 to 3.11 s.

The implementation provides an append-only downstream text contract rather than immutable internal recognition. Internal hypotheses drift in the instrumented subset, tokenizer seams change token segmentation in half the samples, and the limited semantic experiment does not establish equivalence. These findings bound the claim: the architecture substantially weakens the growth of post-speech latency over the tested range while exposing concrete recognition, scheduling, and synthesis trade-offs.

B. Limitations

The main 1,132-sample benchmark remains TTS-synthesized. The added LibriSpeech/AISHELL-1 data are concatenated human-read utterances, not spontaneous dialogue, and some concatenation seams are audible. Babble noise is a demonstrated failure boundary because VAD over-triggering and empty ASR outputs create backlog and long tails. Experiments use exclusive GPUs, a single-machine ASR/LLM queue topology, and localhost TTS; distributed network latency and GPU contention are not represented, and absolute TTFT is CPU- and platform-sensitive.

The downstream interface never rolls back text, but internal re-recognition frequently changes committed-region hypotheses. Fragment-wise tokenization also differs from one-shot tokenization, so KV-state equivalence is not established. Semantic evidence covers only 50 Very Long synthetic samples, independently sampled replies, a 128-token cap, and one

TABLE VIII
UNIFIED TTFA AND COMPONENT DECOMPOSITION ON THE SECOND PLATFORM (MS)

System	Language	Mean $\pm$ std	P50	P90	P95	n
System B	Chinese	3303.3 $\pm$ 2681.3	2603.0	3130.4	8299.2	25
System B	English	7660.5 $\pm$ 2781.2	7577.0	10940.4	11857.2	25
System B	All	5481.9 $\pm$ 3486.1	3113.7	10506.6	11656.3	50
System A	Chinese	22616.8 $\pm$ 2197.2	22161.4	25459.9	26874.8	25
System A	English	22234.7 $\pm$ 2850.4	22439.6	25274.4	26487.5	25
System A	All	22425.7 $\pm$ 2526.1	22269.9	25588.8	26887.4	50
Component means $\pm$ std for all languages						
Component	System B	System A				
Speech end $\rightarrow$ feed end	0.1 $\pm$ 0.0	0.1 $\pm$ 0.0				
Feed end $\rightarrow$ input close	133.0 $\pm$ 98.7	0.0 $\pm$ 0.0				
Input close $\rightarrow$ first token	1383.1 $\pm$ 278.8	4182.9 $\pm$ 898.3				
First token $\rightarrow$ text ready	387.0 $\pm$ 429.2	4681.1 $\pm$ 583.9				
Text ready $\rightarrow$ TTS request	0.4 $\pm$ 0.1	0.6 $\pm$ 0.1				
TTS request $\rightarrow$ playable PCM	3578.3 $\pm$ 3023.6	13561.1 $\pm$ 2513.6				
Closed sum	5481.9	22425.7				

service judge without repeated ratings. The unified TTFA experiment uses sentence-ready System B output versus the non-streaming System A response path, and TTS contributes a large language-dependent tail. Finally, the prototype remains half-duplex, a limitation discussed in work on turn-taking systems [28]. Because the modules communicate through text, the prototype also discards non-textual acoustic cues; this latter point is a direct architectural observation rather than a claim attributed to [28].

C. Future Work

Future work should evaluate spontaneous, accented, microphone, and overlapping speech; add VAD back-pressure and adaptive segmentation for babble; and test task completion with repeated human and multi-judge assessment. Rollback-aware cache repair or retokenization could reconcile later ASR corrections and tokenizer seams. Platform-aware routing should replace a fixed duration threshold. Lower-latency incremental TTS, distributed deployment measurements, GPU-contention tests, larger LLMs, and full-duplex interruption handling would clarify how the observed scheduling gains transfer to production systems.

ACKNOWLEDGMENT

The authors thank the School of Computer, Electronics and Information, Guangxi University, for providing the research environment that supported this work. The authors gratefully acknowledge the open-source projects that made this work possible, including OpenAI Whisper, Silero VAD, and CosyVoice. The authors also thank the anonymous reviewers for their constructive comments.

REFERENCES

- [1] OpenAI, “GPT-4o system card,” *arXiv preprint arXiv:2410.21276*, 2024.
- [2] R. Prabhavalkar, T. Hori, T. N. Sainath, R. Schlüter, and S. Watanabe, “End-to-end speech recognition: a survey,” *IEEE/ACM Transactions on Audio, Speech, and Language Processing*, vol. 32, pp. 325–351, 2024.
- [3] A. Vaswani, N. Shazeer, N. Parmar, J. Uszkoreit, L. Jones, A. N. Gomez, L. Kaiser, and I. Polosukhin, “Attention is all you need,” in *Advances in Neural Information Processing Systems*, vol. 30. Red Hook, NY, USA: Curran Associates, 2017, pp. 5998–6008.

- [4] A. Gulati, J. Qin, C.-C. Chiu, N. Parmar, Y. Zhang, J. Yu, W. Han, S. Wang, Z. Zhang, Y. Wu, and R. Pang, “Conformer: convolution-augmented Transformer for speech recognition,” in *Interspeech 2020*. ISCA, 2020, pp. 5036–5040.
- [5] D. Liu, G. Spanakis, and J. Niehues, “Low-latency sequence-to-sequence speech recognition and translation by partial hypothesis selection,” in *Interspeech 2020*. ISCA, 2020, pp. 3620–3624.
- [6] D. Macháček, R. Dabre, and O. Bojar, “Turning Whisper into real-time transcription system,” in *Proceedings of the 13th International Joint Conference on Natural Language Processing and the 3rd Conference of the Asia-Pacific Chapter of the Association for Computational Linguistics: System Demonstrations*. Association for Computational Linguistics, 2023, pp. 17–24.
- [7] A. Radford, J. W. Kim, T. Xu, G. Brockman, C. McLeavey, and I. Sutskever, “Robust speech recognition via large-scale weak supervision,” in *Proceedings of the 40th International Conference on Machine Learning*, vol. 202. PMLR, 2023, pp. 28 492–28 518.
- [8] Y. Zheng, Y. Chen, B. Qian, X. Shi, Y. Shu, and J. Chen, “A review on edge large language models: design, execution, and applications,” *ACM Computing Surveys*, vol. 57, no. 8, pp. 1–35, 2025, art. no. 209.
- [9] W. Kwon, Z. Li, S. Zhuang, Y. Sheng, L. Zheng, C. H. Yu, J. E. Gonzalez, H. Zhang, and I. Stoica, “Efficient memory management for large language model serving with PagedAttention,” in *Proceedings of the 29th ACM Symposium on Operating Systems Principles*. Koblenz, Germany: ACM, 2023, pp. 611–626.
- [10] T. Dao, D. Y. Fu, S. Ermon, A. Rudra, and C. Ré, “FlashAttention: fast and memory-efficient exact attention with IO-awareness,” in *Advances in Neural Information Processing Systems*, vol. 35. Neural Information Processing Systems Foundation, Inc., 2022, pp. 16 344–16 359.
- [11] G. Xiao, Y. Tian, B. Chen, S. Han, and M. Lewis, “Efficient streaming language models with attention sinks,” in *The Twelfth International Conference on Learning Representations*. Vienna, Austria: OpenReview.net, 2024. [Online]. Available: <https://openreview.net/forum?id=NG7sS51zVF>
- [12] N. Sethiya and C. K. Maurya, “End-to-end speech-to-text translation: a survey,” *Computer Speech & Language*, vol. 90, p. 101751, 2025.
- [13] S. Yin, C. Fu, S. Zhao, K. Li, X. Sun, T. Xu, and E. Chen, “A survey on multimodal large language models,” *National Science Review*, vol. 11, no. 12, p. nwae403, 2024.
- [14] Z. Xie and C. Wu, “Mini-Omni: language models can hear, talk while thinking in streaming,” *arXiv preprint arXiv:2408.16725*, 2024.
- [15] A. Défossez, L. Mazaré, M. Orsini, A. Royer, P. Pérez, H. Jégou, E. Grave, and N. Zeghidour, “Moshi: a speech-text foundation model for real-time dialogue,” *arXiv preprint arXiv:2410.00037*, 2024.
- [16] Q. Fang, S. Guo, Y. Zhou, Z. Ma, S. Zhang, and Y. Feng, “LLaMA-Omni: seamless speech interaction with large language models,” in *The Thirteenth International Conference on Learning Representations*. Singapore: OpenReview.net, 2025. [Online]. Available: <https://openreview.net/forum?id=PYmrUQmMEw>
- [17] OpenAI, “Whisper large-v3-turbo model card and configuration,” Hugging Face model repository, 2024, public

configuration revision 41f01f3fe87f28c78e2fbf8b568835947dd65ed9; accessed August 22, 2026. [Online]. Available: <https://huggingface.co/openai/whisper-large-v3-turbo/tree/41f01f3fe87f28c78e2fbf8b568835947dd65ed9>

- [18] Silero Team, “Silero VAD: pre-trained enterprise-grade voice activity detector,” GitHub repository, 2024, MIT-licensed software; accessed August 22, 2026. [Online]. Available: <https://github.com/snakers4/silero-vad>
- [19] P. Budzianowski, T.-H. Wen, B.-H. Tseng, I. Casanueva, S. Ultes, O. Ramadan, and M. Gašić, “MultiWOZ - a large-scale multi-domain Wizard-of-Oz dataset for task-oriented dialogue modelling,” in *Proceedings of the 2018 Conference on Empirical Methods in Natural Language Processing*. Brussels, Belgium: Association for Computational Linguistics, 2018, pp. 5016–5026.
- [20] Q. Zhu, K. Huang, Z. Zhang, X. Zhu, and M. Huang, “CrossWOZ: a large-scale Chinese cross-domain task-oriented dialogue dataset,” *Transactions of the Association for Computational Linguistics*, vol. 8, pp. 281–295, 2020.
- [21] Z. Du, Q. Chen, S. Zhang, K. Hu, H. Lu, Y. Yang, H. Hu, S. Zheng, Y. Gu, Z. Ma, Z. Gao, and Z. Yan, “CosyVoice: a scalable multilingual zero-shot text-to-speech synthesizer based on supervised semantic tokens,” *arXiv preprint arXiv:2407.05407*, 2024.
- [22] V. Panayotov, G. Chen, D. Povey, and S. Khudanpur, “LibriSpeech: an ASR corpus based on public domain audio books,” in *2015 IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP)*. IEEE, 2015, pp. 5206–5210.
- [23] H. Bu, J. Du, X. Na, B. Wu, and H. Zheng, “AISHELL-1: an open-source Mandarin speech corpus and a speech recognition baseline,” in *2017 20th Conference of the Oriental Chapter of the International Coordinating Committee on Speech Databases and Speech I/O Systems and Assessment (O-COCOSDA)*. IEEE, 2017, pp. 1–5.
- [24] D. Snyder, G. Chen, and D. Povey, “MUSAN: a music, speech, and noise corpus,” *arXiv preprint arXiv:1510.08484*, 2015.
- [25] A. Yang, B. Yang, B. Hui, B. Zheng, B. Yu, C. Zhou *et al.*, “Qwen2 technical report,” *arXiv preprint arXiv:2407.10671*, 2024.
- [26] J. Chen, S. Xiao, P. Zhang, K. Luo, D. Lian, and Z. Liu, “M3-Embedding: multi-linguality, multi-functionality, multi-granularity text embeddings through self-knowledge distillation,” in *Findings of the Association for Computational Linguistics: ACL 2024*. Association for Computational Linguistics, 2024, pp. 2318–2335.
- [27] Command Code, “DeepSeek V4 Flash (latest): pricing, benchmarks, and speed,” Command Code online model catalog and service documentation, Jul. 2026, service model identifier: deepseek/deepseek-v4-flash; page release date July 31, 2026; accessed August 22, 2026. [Online]. Available: <https://commandcode.ai/models/deepseek-v4-flash>
- [28] G. Skantze, “Turn-taking in conversational systems and human-robot interaction: a review,” *Computer Speech & Language*, vol. 67, p. 101178, 2021.