




*The Conv Stem reduces temporal resolution before Transformer encoding; the decoder generates text and timestamp tokens via cross-attention.*