

Total Input Window to Model

Transcription Target

Prefix Context

(frozen history · not re-emitted)

Current Segment(s)

(active inference → emitted)

Suffix Buffer

(future context · protects unstable tail)

incoming audio stream

t_{start}

t_{current}

t_{end}

Time t